


Deep Learning Basics

Lecture 5: Modern Convolutional Neural Networks

최성준 (고려대학교 인공지능학과)

AlexNet



Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, "ImageNet Classification with Deep Convolutional Neural Networks," NIPS, 2012

ILSVRC

- ImageNet Large-Scale Visual Recognition Challenge

- Classification / Detection / Localization / Segmentation

- 1,000 different categories

- Over 1 million images

- Training set: 456,567 images

Image classification

Easiest classes



Hardest classes




ILSVRC



Year	Error Rate
2010	28.2%
2011	25.8%
2012	16.4%
2013	11.2%
2014	6.7%
2015	3.5%
Human	About 5.1%

AlexNet




AlexNet

- Key ideas
 - Rectified Linear Unit (ReLU) activation
 - GPU implementation (2 GPUs)
 - Local response normalization, Overlapping pooling
 - Data augmentation
 - Dropout


AlexNet

- ReLU Activation
 - Preserves properties of linear models
 - Easy to optimize with gradient descent
 - Good generalization
 - Overcome the vanishing gradient problem




VGGNet

Karen Simonyan, Andrew Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," ICLR, 2015




VGGNet



- Increasing depth with 3×3 convolution filters (with stride 1)
- 1x1 convolution for fully connected layers
- Dropout ($p=0.5$)
- VGG16, VGG19

VGGNet

- Why 3×3 convolution?




128

3×3

128

3×3

128



128

5×5

128

Receptive field

5×5

5×5


of params

$$3 \times 3 \times 128 \times 128 + 3 \times 3 \times 128 \times 128 = \mathbf{294,912}$$

$$5 \times 5 \times 128 \times 128 = \mathbf{409,600}$$

GoogLeNet

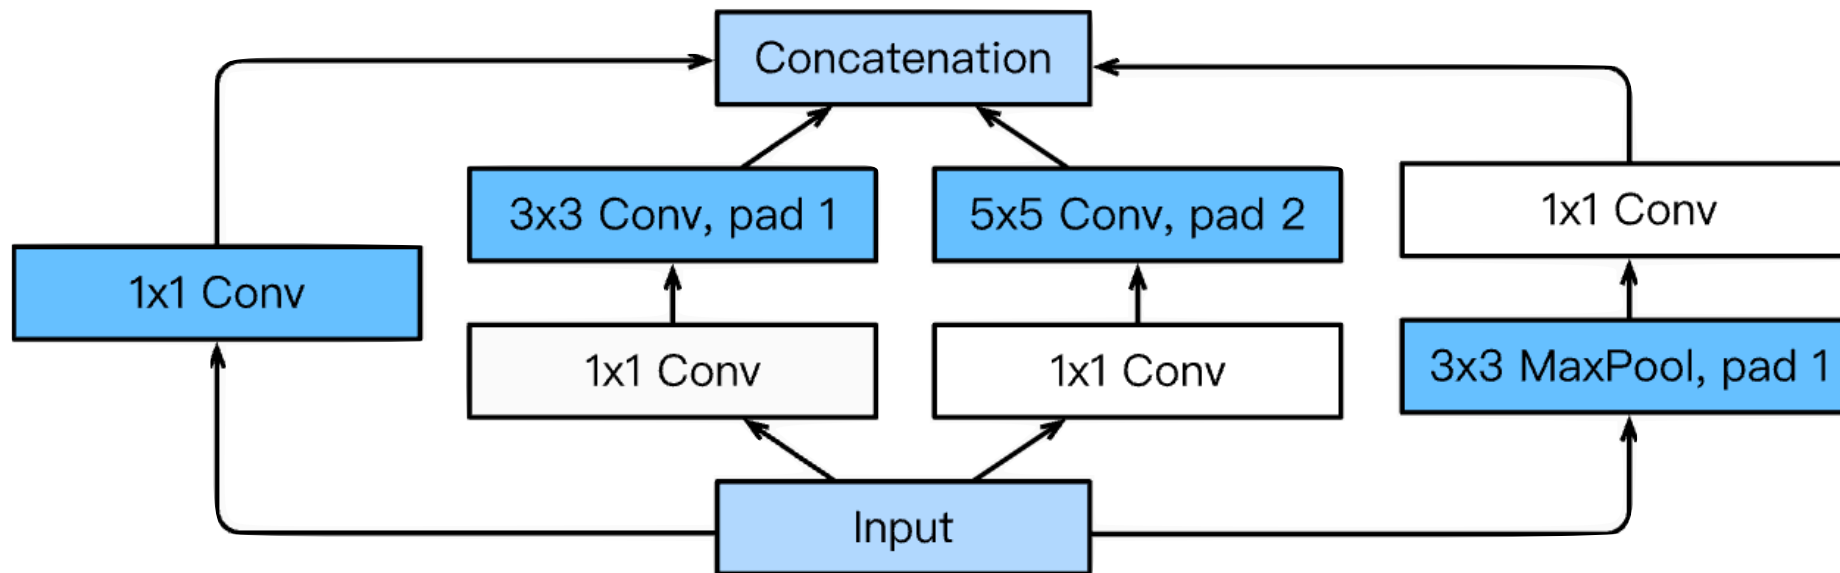
Christian et al. “Going Deeper with Convolutions”, CVPR, 2015



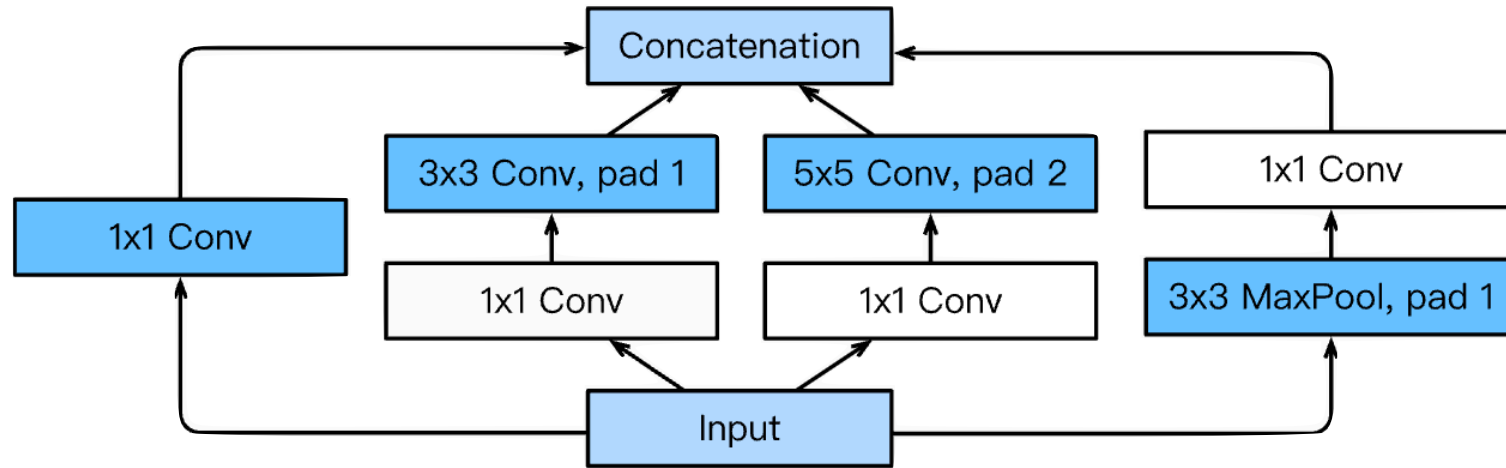


GoogLeNet

- GoogLeNet won the ILSVRC at 2014
 - It combined network-in-network (NiN) with inception blocks.
- Inception blocks




Inception Block




- What are the benefits of the inception block?
 - Reduce the number of parameter.
- How?
 - Recall how the number of parameters is computed.
 - 1x1 convolution can be seen as channel-wise dimension reduction.

Inception Block

- Benefit of 1x1 convolution



$$3 \times 3 \times 128 \times 128 = 147,456$$



$$1 \times 1 \times 128 \times 32 = 4,096$$

$$3 \times 3 \times 32 \times 128 = 36,864$$

$$4,096 + 36,864 = 40,960$$

1x1 convolution enables about 30% reduce of the number of parameters!

Quiz

● Which CNN architecture has the least number of parameters?


1. AlexNet (8-layers) (60M)

2. VGGNet (19-layers) (110M)

3. GoogLeNet (22-layers) (4M)

● The answer is **GoogLeNet**.


ResNet



Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun, "Deep Residual Learning for Image Recognition," CVPR, 2015

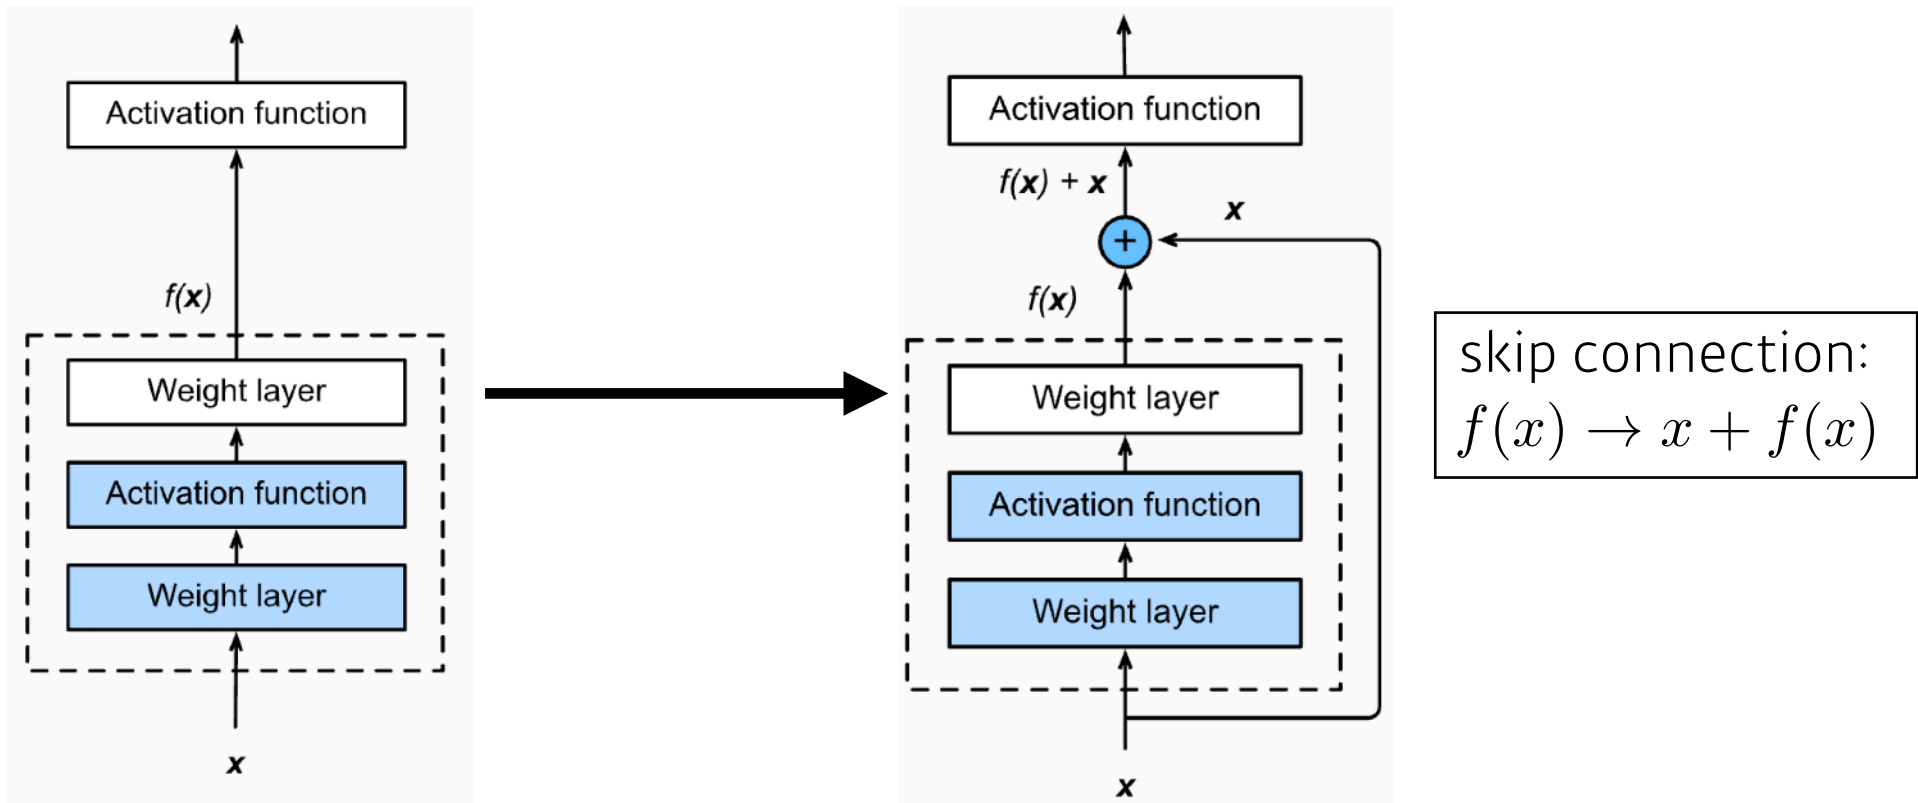
ResNet

- Deeper neural networks are hard to train.
 - Overfitting is usually caused by an excessive number of parameters.
 - But, not in this case.





ResNet

- Add an identity map (skip connection)




ResNet

- Add an identity map (skip connection)




ResNet

- Add an identity map **after** nonlinear activations:



Simple Shortcut




Projected Shortcut


1x1 convolution to match the channel depth

ResNet

- Batch normalization **after** convolutions:



Simple Shortcut





Projected Shortcut



1x1 convolution to match the channel depth

ResNet

- Bottleneck architecture



ResNet




● **Performance** increases while **parameter size** decreases.

DenseNet


Gao Huang, Zhuang Liu, Laurens van der Maaten, Kilian Weinberger,
“Densely Connected Convolutional Networks,” CVPR, 2017

DenseNet

- DenseNet uses **concatenation** instead of **addition**.




ResNet



DenseNet

DenseNet

- DenseNet uses **concatenation** instead of **addition**.



$$\mathbf{x} \mapsto [\mathbf{x}, f_1(\mathbf{x}), f_2(\mathbf{x}, f_1(\mathbf{x})), f_3(\mathbf{x}, f_1(\mathbf{x}), f_2(\mathbf{x}, f_1(\mathbf{x}))), f_4(\mathbf{x}, f_1(\mathbf{x}), f_2(\mathbf{x}, f_1(\mathbf{x})), f_3(\mathbf{x}, f_1(\mathbf{x}), f_2(\mathbf{x}, f_1(\mathbf{x})))]$$

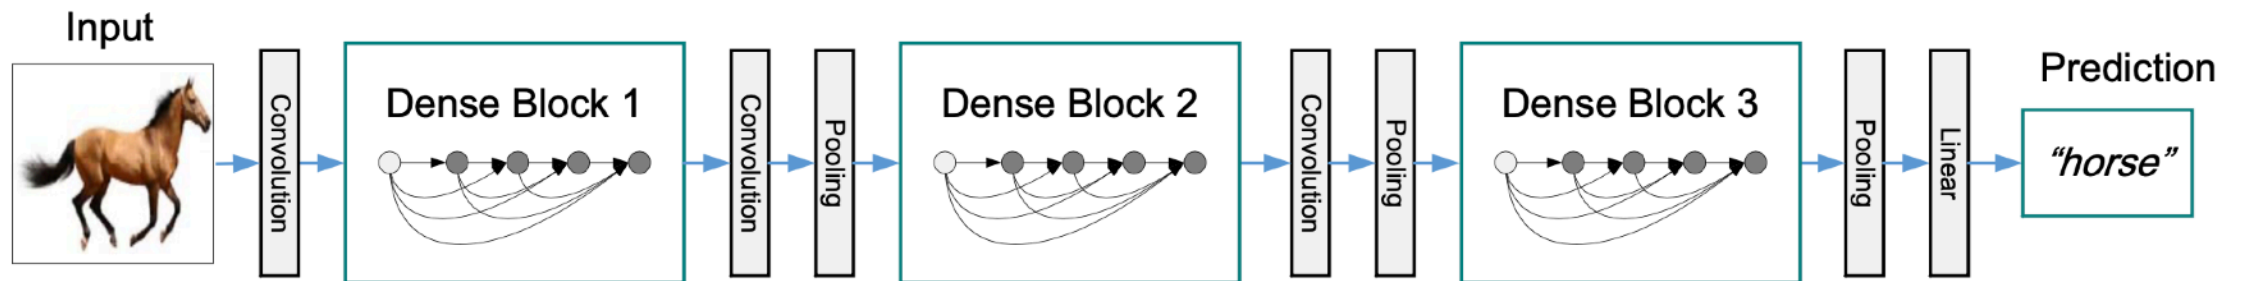
DenseNet

● Dense Block

- Each layer concatenates the feature maps of all preceding layers.
- The number of channels increases geometrically.

● Transition Block

- BatchNorm -> 1x1 Conv -> 2x2 AvgPooling
- Dimension reduction



Summary

- Key takeaways
 - **VGG**: repeated 3x3 blocks
 - **GoogLeNet**: 1x1 convolution
 - **ResNet**: skip-connection
 - **DenseNet**: concatenation

Thank you for listening
