

자연어 처리 DAY 2

Jaegul Choo

Associate Professor, Graduate School of AI, KAIST


1.

Basics of Recurrent Neural Networks (RNNs)

Basic problem settings

Model architecture and how it works


- Basic structure



An unrolled recurrent neural network.

<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- Inputs and outputs of RNNs (rolled version)
 - We usually want to predict a vector at some time steps




http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf

- **How to calculate the hidden state of RNNs**

- We can process a sequence of vectors by applying a recurrence formula at every time step
- h_{t-1} : old hidden-state vector
- x_t : input vector at some time step
- h_t : new hidden-state vector
- f_W : RNN function with parameters W
- y_t : output vector at time step t


$$h_t = f_W(h_{t-1}, x_t)$$



http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf


- How to calculate the hidden state of RNNs
 - Notice: The same function and the same set of parameters are used at every time step

$$h_t = f_W(h_{t-1}, x_t)$$



http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf

- How to calculate the hidden state of RNNs
 - The state consists of a single “hidden” vector \underline{h}



$$h_t = f_W(h_{t-1}, x_t)$$



$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

End of Document
Thank You.

2.

Types of RNNs

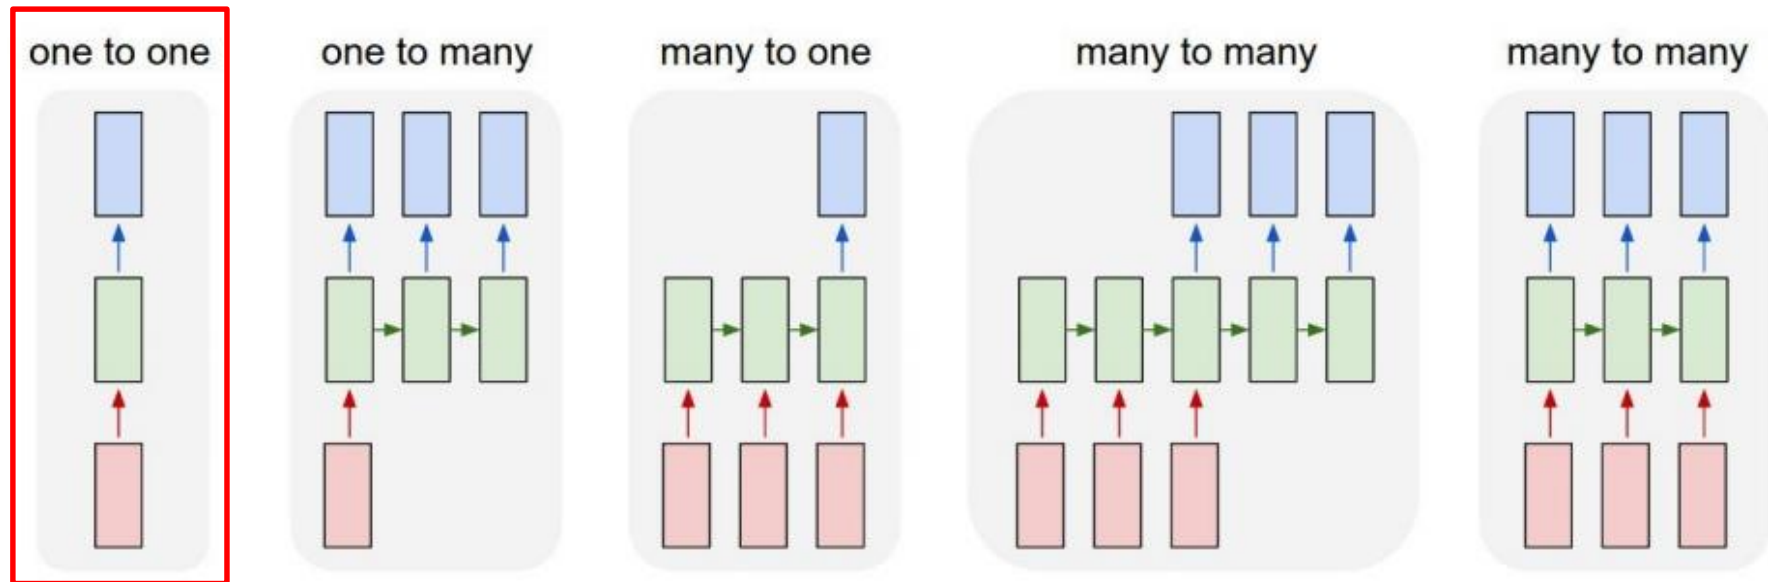
one-to-one

one-to-many

many-to-one


many-to-many

- One-to-one
 - Standard Neural Networks



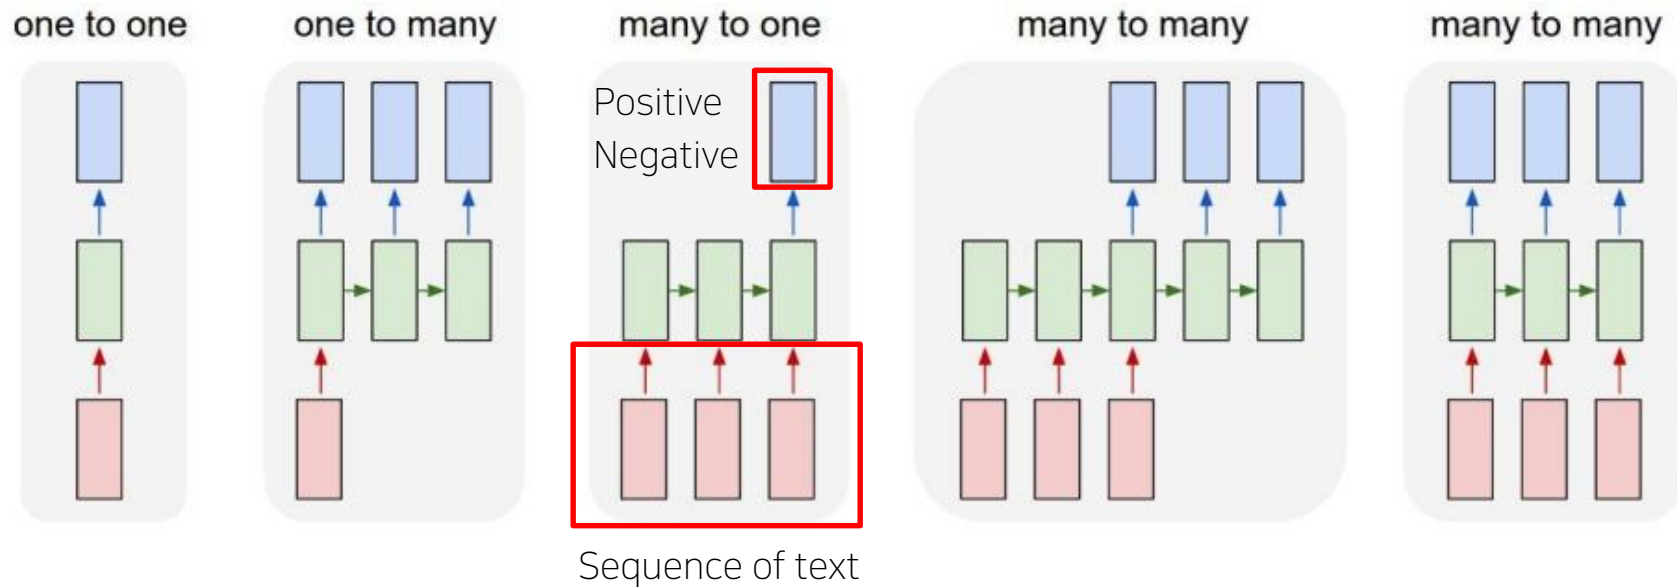
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- One-to-many
 - Image Captioning




<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Many-to-one
 - Sentiment Classification




<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Sequence-to-sequence
 - Machine Translation



<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- **Sequence-to-sequence**
 - Video classification on frame level




<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

End of Document
Thank You.

3.


Character-level Language Model

- Example of training sequence "hello"
 - Vocabulary: [h, e, l, o]
 - Example training sequence: "hello"




<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- **Example of training sequence "hello"**
 - Vocabulary: [h, e, l, o]
 - Example training sequence: "hello"



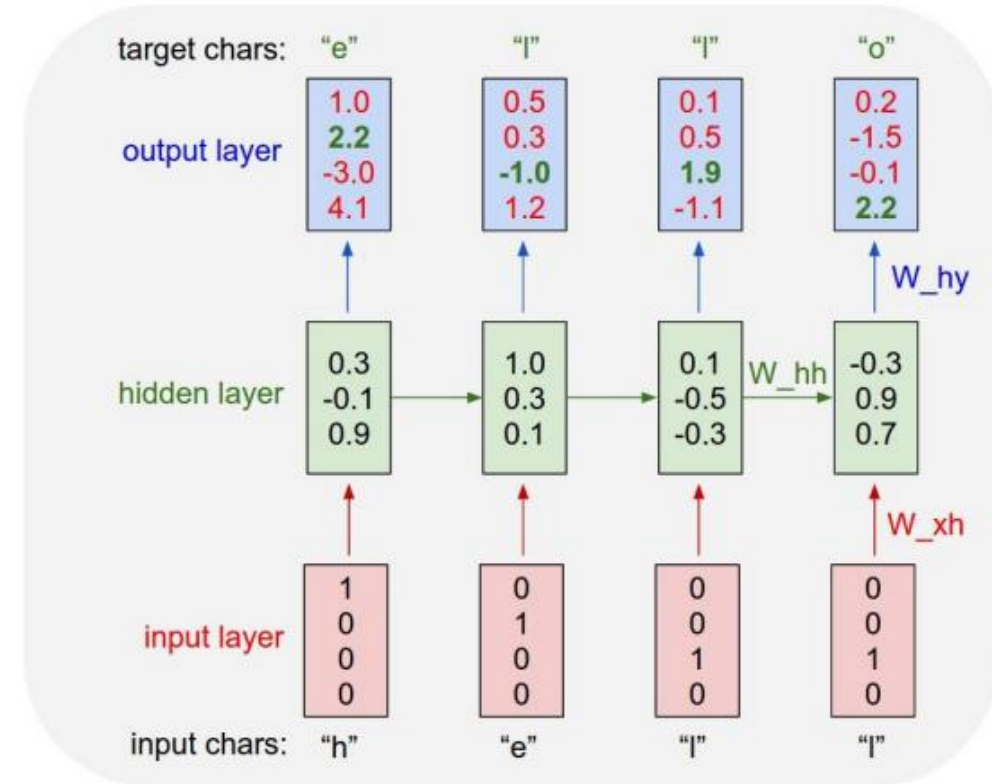
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Example of training sequence "hello"
 - $h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t + b)$



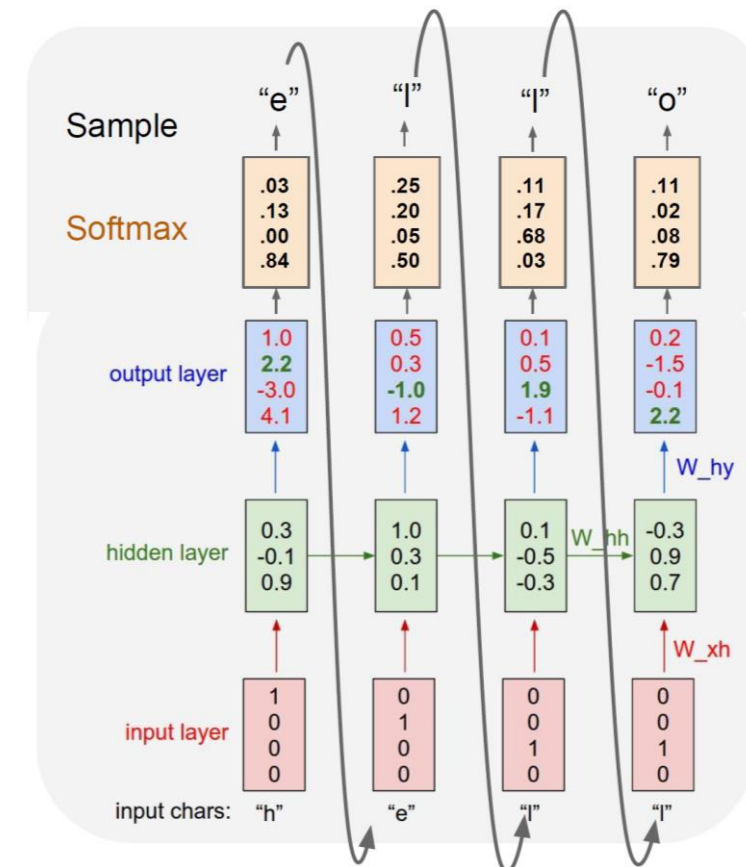
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Example of training sequence "hello"
 - $\text{Logit} = W_{hy}h_t + b$



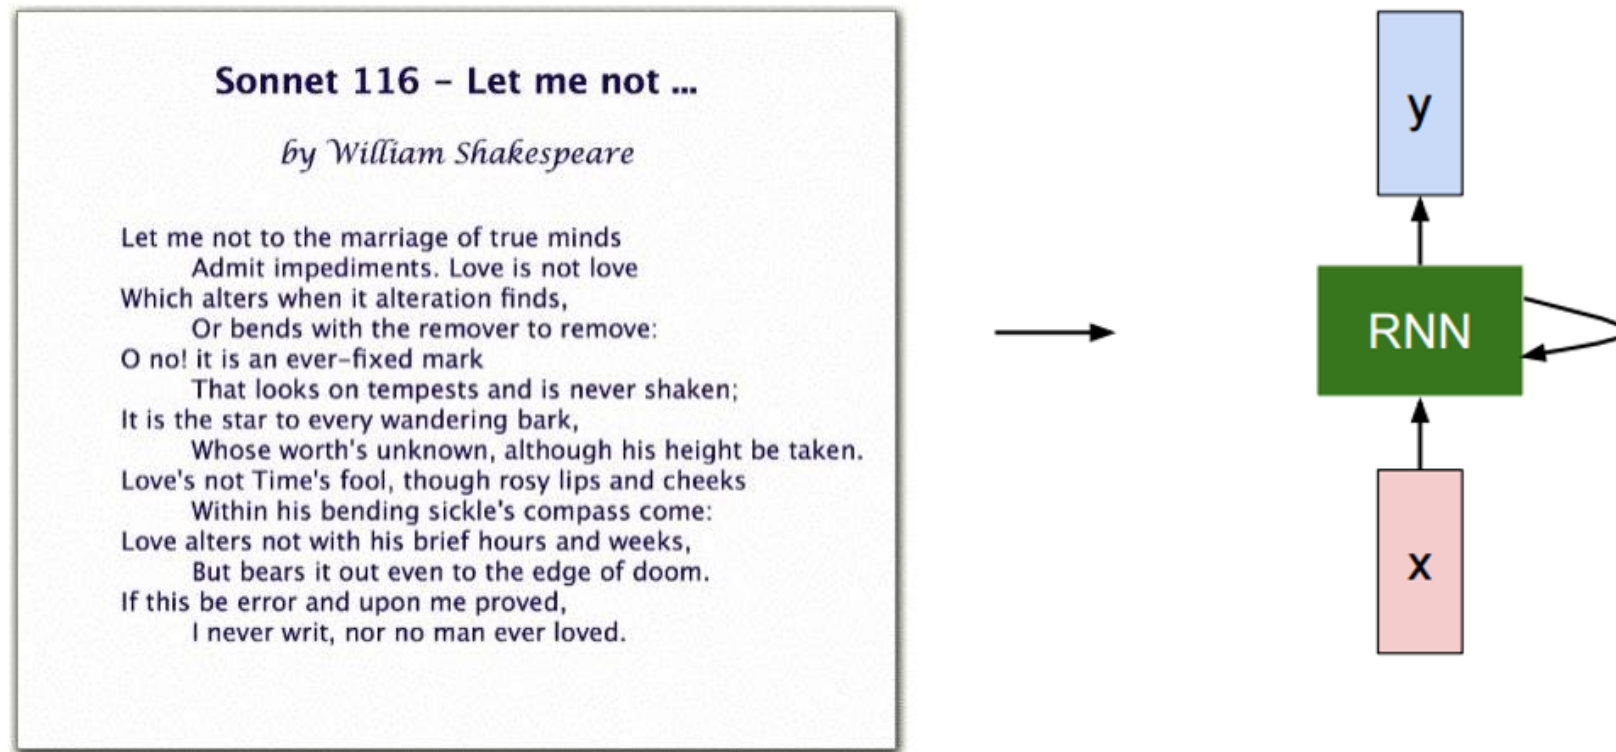
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Example of training sequence "hello"
 - At test-time, sample characters one at a time, feed back to model



<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Training a RNN on Shakespeare's plays



http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf

- Training process of RNN

at first:

```
tyntd-iafhatawiaoihrdemot lytdws e ,tfti, astai f ogoh eoase rrranbyne 'nhthnee e  
plia tklrkd t o idoe ns,smtt h ne etie h,hregtrs nigtkie,aoaenns lng
```

train more

```
"Tmont thithey" fomesscerliund  
Keushey. Thom here  
sheulke, anmerenith ol sivh I lalterthend Bleipile shuwu fil on aseterlome  
coaniogennc Phe lism thond hon at. MeiDimorotion in ther thize."
```

train more

```
Aftair fall unsuch that the hall for Prince Velzonski's that me of  
her hearly, and behs to so arwage fiving were to it beloge, pavu say falling misfort  
how, and Gogition is so overelical and offer.
```

train more

```
"Why do what that day," replied Natasha, and wishing to himself the fact the  
princess, Princess Mary was easier, fed in had oftended him.  
Pierre aking his soul came to the packs and drove up his father-in-law women.
```

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- Results of trained RNN

PANDARUS:

Alas, I think he shall be come approached and the day
When little strain would be attain'd into being never fed,
And who is but a chain and subjects of his death,
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,
Breaking and strongly should be buried, when I perish
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and
my fair nudes begun out of the fact, to be conveyed,
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.

VIOLA:

Why, Salisbury must find his flesh and thought
That which I am not apt, not a man and in fire,
To show the reining of the raven and the wars
To grace my hand reproach within, and not a fair are hand,
That Caesar and my goodly father's world;
When I was heaven of presence and our fleets,
We spare with hours, but cut thy council I am great,
Murdered and by thy master's ready there
My power to give thee but so much as hell:
Some service in the noble bondman here,
Would show him to her wine.

KING LEAR:

O, if you were a feeble sight, the courtesy of your law,
Your sight and several breath, will wear the gods
With his heads, and my hands are wonder'd at the deeds,
So drop upon your lordship's head, and your opinion
Shall be against your honour.

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- A paper written by RNN

Proof. Omitted. □

Lemma 0.1. *Let \mathcal{C} be a set of the construction.*
Let \mathcal{C} be a gerber covering. Let \mathcal{F} be a quasi-coherent sheaves of \mathcal{O} -modules. We have to show that

$$\mathcal{O}_{\mathcal{O}_X} = \mathcal{O}_X(\mathcal{L})$$

Proof. This is an algebraic space with the composition of sheaves \mathcal{F} on $X_{\acute{e}tale}$ we have

$$\mathcal{O}_X(\mathcal{F}) = \{morph_1 \times_{\mathcal{O}_X} (\mathcal{G}, \mathcal{F})\}$$

where \mathcal{G} defines an isomorphism $\mathcal{F} \rightarrow \mathcal{F}$ of \mathcal{O} -modules. □

Lemma 0.2. *This is an integer \mathbb{Z} is injective.*
Proof. See Spaces, Lemma ?? □

Lemma 0.3. *Let S be a scheme. Let X be a scheme and X is an affine open covering. Let $\mathcal{U} \subset \mathcal{X}$ be a canonical and locally of finite type. Let X be a scheme. Let X be a scheme which is equal to the formal complex.*
The following to the construction of the lemma follows.
Let X be a scheme. Let X be a scheme covering. Let

$$b : X \rightarrow Y' \rightarrow Y \rightarrow Y' \times_X Y \rightarrow X.$$

be a morphism of algebraic spaces over S and Y .
Proof. Let X be a nonzero scheme of X . Let X be an algebraic space. Let \mathcal{F} be a quasi-coherent sheaf of \mathcal{O}_X -modules. The following are equivalent

- (1) \mathcal{F} is an algebraic space over S .
- (2) If X is an affine open covering.

Consider a common structure on X and X the functor $\mathcal{O}_X(\mathcal{U})$ which is locally of finite type. □

This since $\mathcal{F} \in \mathcal{F}$ and $x \in \mathcal{G}$ the diagram

is a limit. Then \mathcal{G} is a finite type and assume S is a flat and \mathcal{F} and \mathcal{G} is a finite type f_* . This is of finite type diagrams, and

- the composition of \mathcal{G} is a regular sequence,
- $\mathcal{O}_{X'}$ is a sheaf of rings.

□

Proof. We have see that $X = \text{Spec}(R)$ and \mathcal{F} is a finite type representable by algebraic space. The property \mathcal{F} is a finite morphism of algebraic stacks. Then the cohomology of X is an open neighbourhood of U . □

Proof. This is clear that \mathcal{G} is a finite presentation, see Lemmas ??.
A reduced above we conclude that U is an open covering of \mathcal{C} . The functor \mathcal{F} is a "field"

$$\mathcal{O}_{X,x} \rightarrow \mathcal{F}_x \rightarrow \mathcal{O}_{X_{\acute{e}tale}} \rightarrow \mathcal{O}_{X_{\acute{e}tale}}^{-1} \mathcal{O}_{X_{\acute{e}tale}}(\mathcal{O}_{X_{\acute{e}tale}}^{\vee})$$

is an isomorphism of covering of $\mathcal{O}_{X_{\acute{e}tale}}$. If \mathcal{F} is the unique element of \mathcal{F} such that X is an isomorphism.
 The property \mathcal{F} is a disjoint union of Proposition ?? and we can filtered set of presentations of a scheme \mathcal{O}_X -algebra with \mathcal{F} are opens of finite type over S .
 If \mathcal{F} is a scheme theoretic image points. □

If \mathcal{F} is a finite direct sum $\mathcal{O}_{X_{\acute{e}tale}}$ is a closed immersion, see Lemma ?? . This is a sequence of \mathcal{F} is a similar morphism.

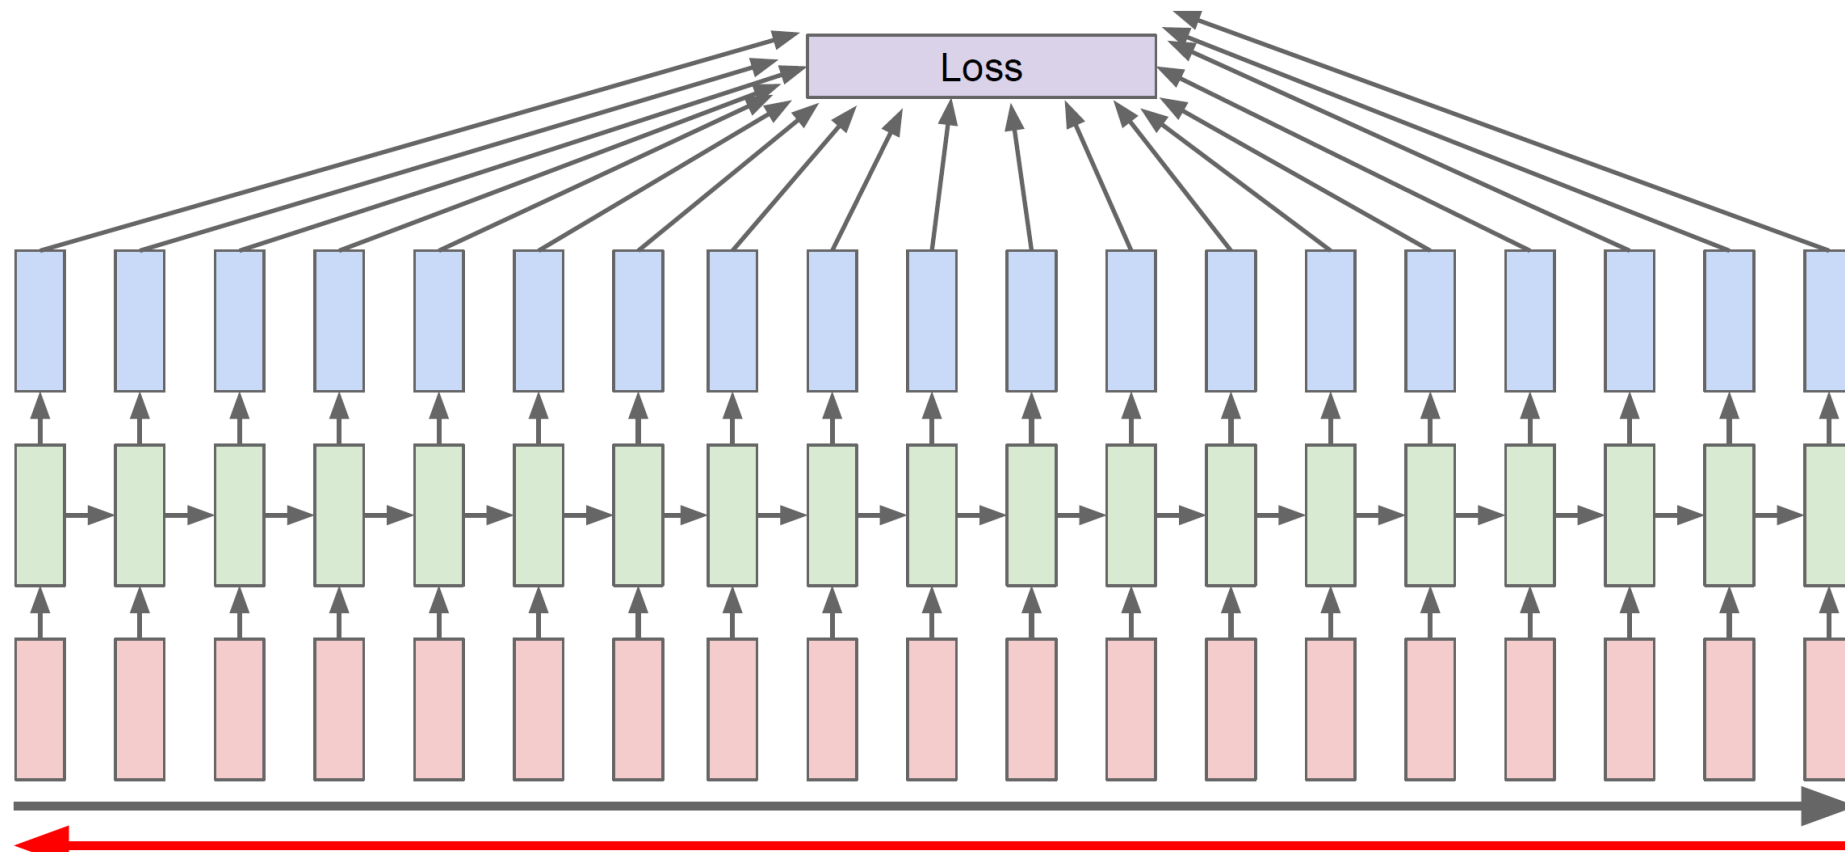
- C code generated by RNN

```
static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << 1))
            pipe = (in_use & UMXTHREAD_UNCCA) +
                ((count & 0x00000000ffffffff) & 0x0000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &soffset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(s, "policy ");
}
```

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Backpropagation through time (BPTT)


- Forward through entire sequence to compute loss, then backward through entire sequence to compute gradient



Backpropagation through time (BPTT)

Character-level Language Model

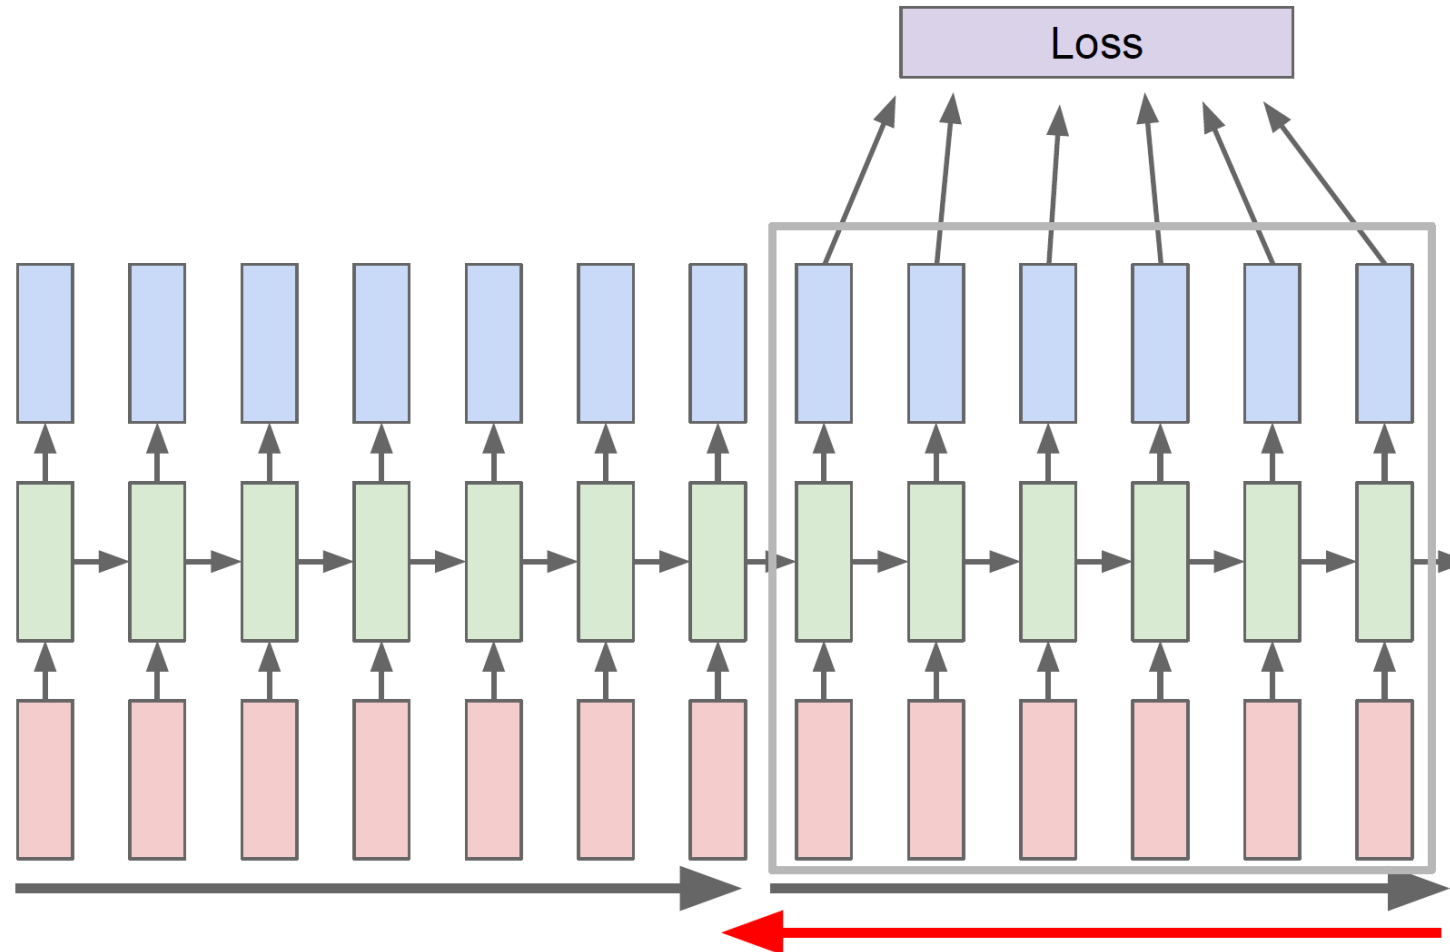
- Run forward and backward through chunks of the sequence instead of whole sequence



Backpropagation through time (BPTT)


Character-level Language Model

- Carry hidden states forward in time forever, but only backpropagate for some smaller number of steps



Backpropagation through time (BPTT)


Character-level Language Model



http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf

Searching for Interpretable Cells

Character-level Language Model



http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf

- How RNN works

```
/* Unpack a filter field's string representation from user-space
 * buffer. */
char *audit_unpack_string(void **bufp, size_t *remain, size_t len)
{
    char *str;
    if (!*bufp || (len == 0) || (len > *remain))
        return ERR_PTR(-EINVAL);
    /* of the currently implemented string fields, PATH_MAX
     * defines the longest valid length.
     */
}
```

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- How RNN works
 - Quote detection cell

"You mean to imply that I have nothing to eat out of.... On the contrary, I can supply you with everything even if you want to give dinner parties," warmly replied Chichagov, who tried by every word he spoke to prove his own rectitude and therefore imagined Kutuzov to be animated by the same desire.

Kutuzov, shrugging his shoulders, replied with his subtle penetrating smile: "I meant merely to say what I said."


<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- How RNN works
 - If statement cell

```
static int __dequeue_signal(struct sigpending *pending, sigset_t *mask,
                           siginfo_t *info)
{
    int sig = next_signal(pending, mask);
    if (sig) {
        if (current->notifier) {
            if (sigismember(current->notifier_mask, sig)) {
                if (!(current->notifier)(current->notifier_data)) {
                    clear_thread_flag(TIF_SIGPENDING);
                    return 0;
                }
            }
        }
        collect_signal(sig, pending, info);
    }
    return sig;
}
```

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

- **RNN is excellent, but...**
 - Multiplying the same matrix at each time step during backpropagation causes gradient vanishing or exploding



- Toy Example

- $h_t = \tanh(w_{xh}x_t + w_{hh}h_{t-1} + b), t = 1, 2, 3$
- For $w_{hh} = 3, w_{xh} = 2, b = 1$

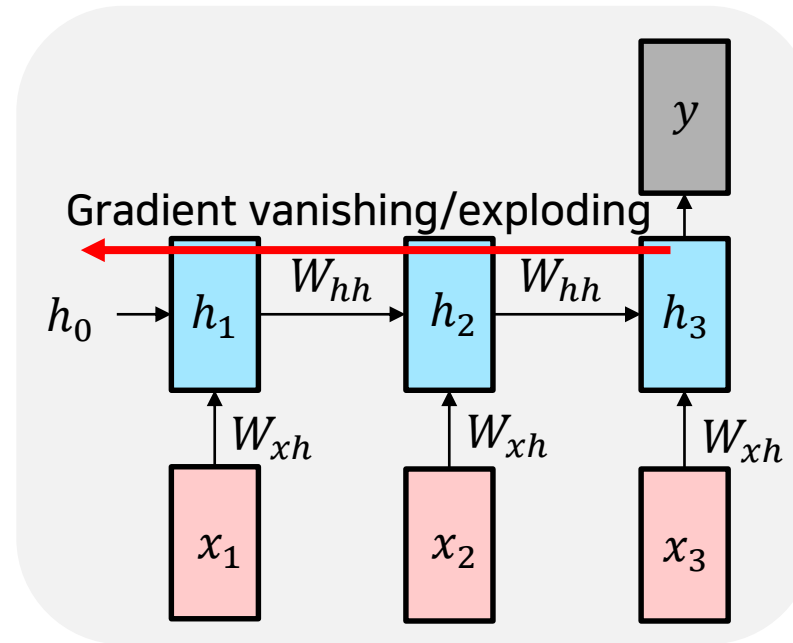
$$h_3 = \tanh(2x_3 + 3h_2 + 1)$$

$$h_2 = \tanh(2x_2 + 3h_1 + 1)$$


$$h_1 = \tanh(2x_1 + 3h_0 + 1)$$

...

$$h_3 = \tanh(2x_3 + 3 \tanh(2x_2 + 3 \tanh(2x_1 + 3h_0 + 1) + 1) + 1)$$



- The reason why the vanishing gradient problem is important




<https://imgur.com/gallery/vaNahKE>

End of Document
Thank You.

4.

Long Short-Term Memory (LSTM)
Gated Recurrent Unit (GRU)


- **Core Idea:** pass cell state information straightly without any transformation
 - Solving long-term dependency problem



The repeating module in an LSTM contains four interacting layers.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>


- What is LSTM (Long Short-Term Memory)?



The repeating module in an LSTM contains four interacting layers.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- Long short-term memory
 - i: Input gate, Whether to write to cell
 - f: Forget gate, Whether to erase cell
 - o: Output gate, How much to reveal cell
 - g: Gate gate, How much to write to cell




$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$
$$h_t = o \odot \tanh(c_t)$$

Long Short-Term Memory (LSTM)

LSTM, GRU


- A gate exists for controlling how much information could flow from cell state



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>


- Forget gate

- $f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$




<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- Generate information to be added and cut it by input gate
 - $i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$
 - $\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$
- Generate new cell state by adding current information to previous cell state
 - $C_t = f_t \cdot C_{t-1} + i_t \cdot \tilde{C}_t$



[/colah.github.io/posts/2015-08-Understanding-LSTMs/](https://colah.github.io/posts/2015-08-Understanding-LSTMs/)


- Generate hidden state by passing cell state to tanh and output gate
- Pass this hidden state to next time step, and output or next layer if needed
 - $o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$
 - $h_t = o_t \cdot \tanh(C_t)$



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- What is GRU?

- $z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$
- $r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$
- $\tilde{h}_t = \tanh(W \cdot [r_t \cdot h_{t-1}, x_t])$
- $h_t = (1 - z_t) \cdot h_{t-1} + z_t \cdot \tilde{h}_t$
- c.f) $C_t = f_t \cdot C_{t-1} + i_t \cdot \tilde{C}_t$
in LSTM

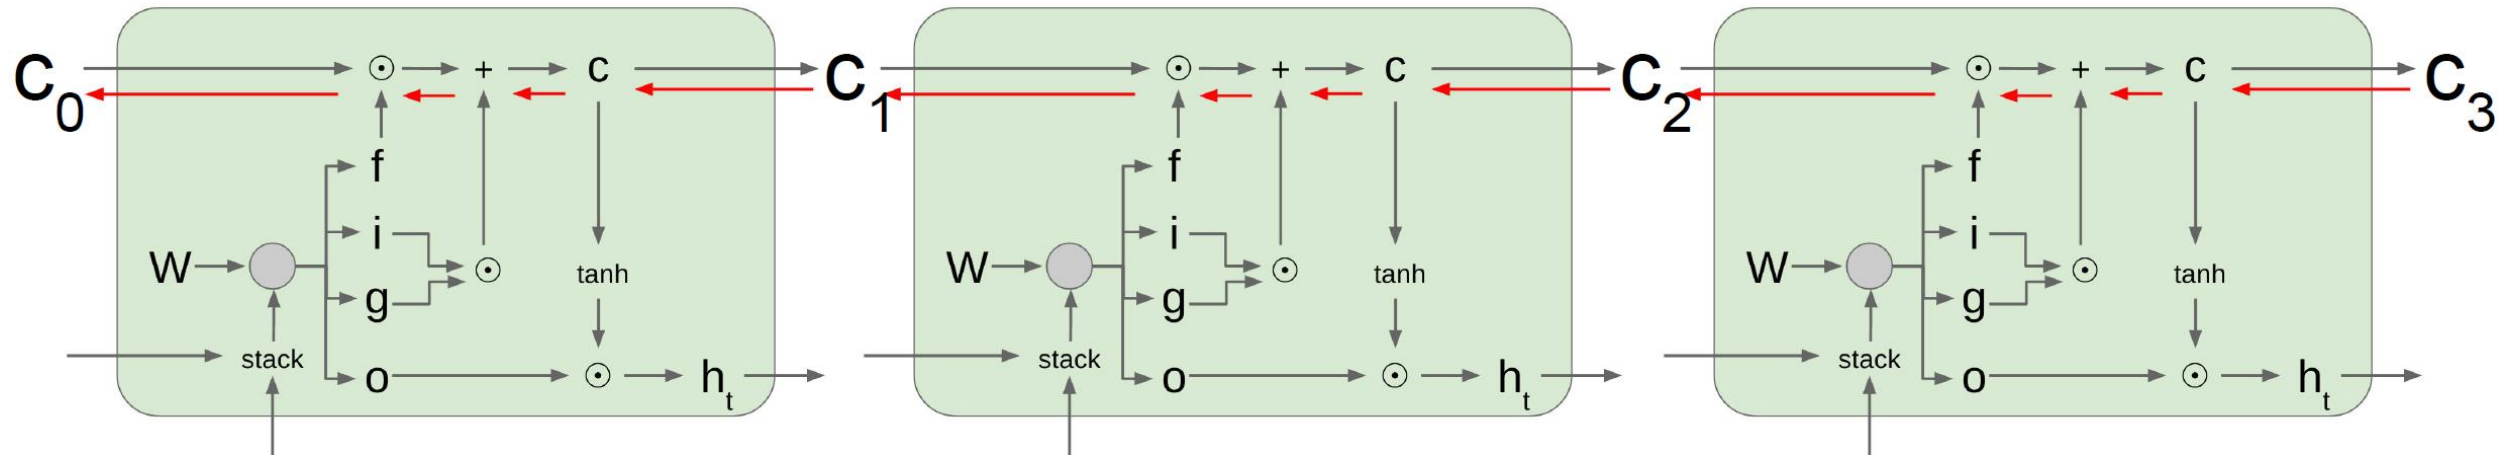


<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Backpropagation in LSTM?GRU

LSTM, GRU

- Uninterrupted gradient flow !



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

- RNNs allow a lot of **flexibility** in architecture design
- Vanilla RNNs are **simple** but don't work very well
- Backward flow of gradients in RNN can **explode or vanish**
- Common to use LSTM or GRU: their additive interactions **improve gradient flow**

References

- <https://imgur.com/gallery/vaNahKE>
- <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>
- <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf

End of Document
Thank You.